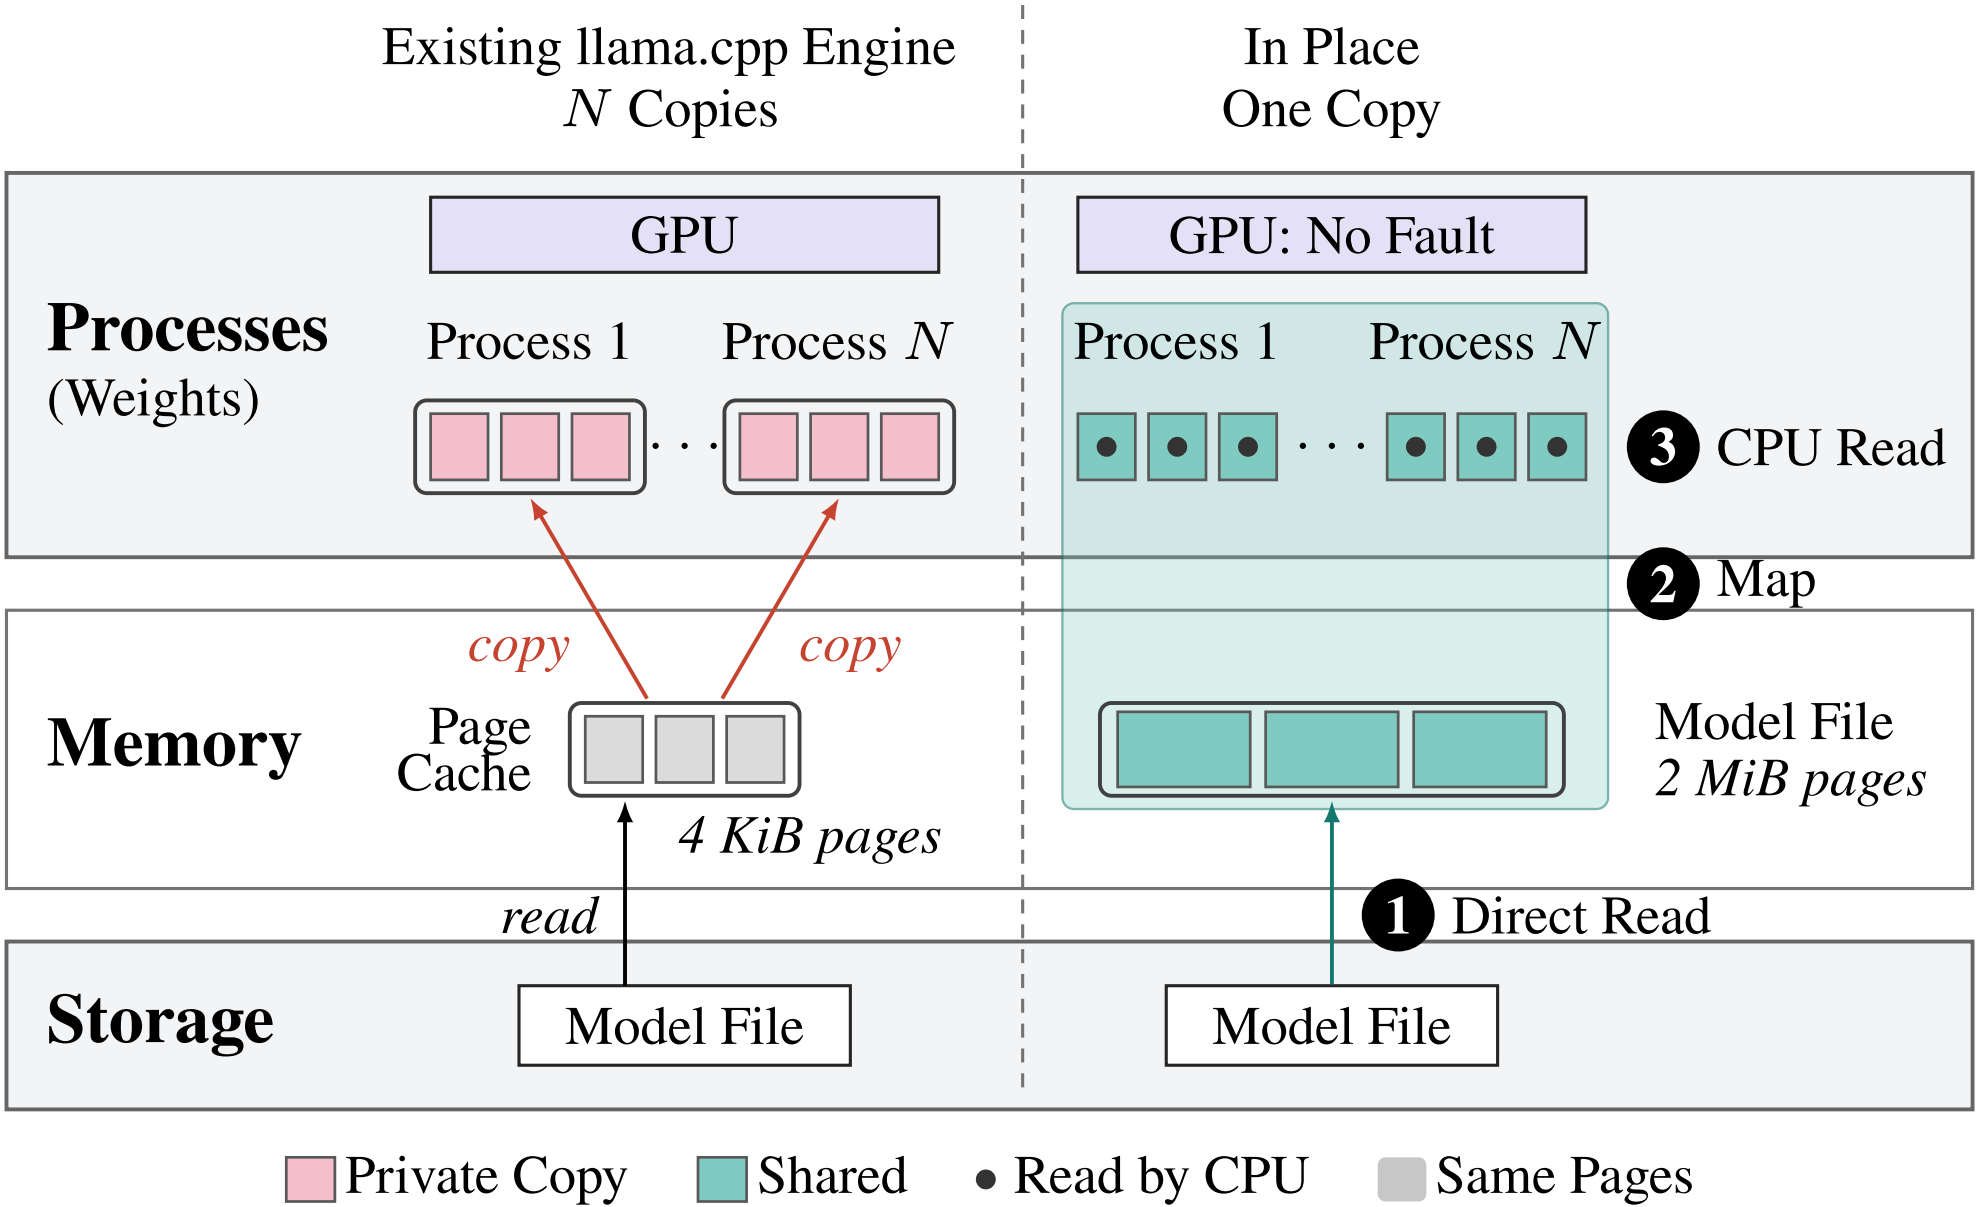

Existing llama.cpp Engine
N Copies
In Place
One Copy
Processes
(Weights)
GPU
GPU: No Fault
Process 1
Process N
Process 1
Process N
3 CPU Read
2 Map
Memory
copy
copy
Page
Cache
4 KiB pages
Model File
2 MiB pages
1 Direct Read
read
Storage
Model File
Model File
Private Copy
Shared
Read by CPU
Same Pages