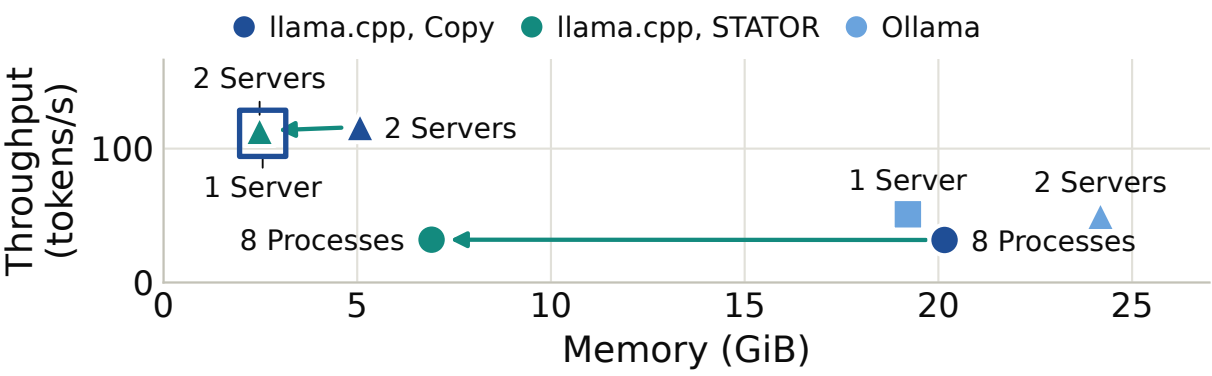
llama.cpp, Copy
llama.cpp, STATOR
Ollama
Throughput (tokens/s)
100
0
2 Servers
2 Servers
1 Server
1 Server
2 Servers
8 Processes
8 Processes
0
5
10
15
20
25
Memory (GiB)